

Homework 2

20200639 Chae Woo Jin

Chapter 4: 1, 4, 10, 12, 14

Chapter 5: 1, 2, 5, 6, 7

Chapter 4

#1.

$$(4.2) \quad p(X) = \frac{e^{\beta_0 + \beta_1 X}}{1 + e^{\beta_0 + \beta_1 X}}, \quad \frac{p(X)}{1 - p(X)} = \frac{\frac{e^{\beta_0 + \beta_1 X}}{1 + e^{\beta_0 + \beta_1 X}}}{\frac{1}{1 + e^{\beta_0 + \beta_1 X}}} = e^{\beta_0 + \beta_1 X} \quad (4.3)$$

#4.

(a) On average, $\frac{1}{10}$ will be used to make the prediction, because observations are uniformly distributed.

(b) On average, when $p = 2$, $\frac{1}{10} \times \frac{1}{10} = \frac{1}{100}$ will be used.

(c) On average, when $p = 100$, $\left(\frac{1}{10}\right)^{100}$ will be used.

(d) As p increases linearly, the proportion of observations against the whole data set which are used to predict the real value decreases exponentially.

(e) The length of p -dimensional hypercube which contains 10% of the training observation is $\sqrt[p]{\frac{1}{10}}$, so that the volume of this hypercube is $\frac{1}{10}$, which is a proportion that it contains its training observations.

#10.

$$\begin{aligned}
\log\left(\frac{\Pr(Y = k|X = x)}{\Pr(Y = K|X = x)}\right) &= \log\left(\frac{\pi_k f_k(x)}{\pi_K f_K(x)}\right) \\
&= \log\left(\frac{\pi_k \exp\left(-\frac{1}{2}(x - \mu_k)^2/\sigma^2\right)}{\pi_K \exp\left(-\frac{1}{2}(x - \mu_K)^2/\sigma^2\right)}\right) \\
&= \log\left(\frac{\pi_k}{\pi_K}\right) - \frac{1}{2} \frac{(x - \mu_k)^2}{\sigma^2} + \frac{1}{2} \frac{(x - \mu_K)^2}{\sigma^2} \\
&= \log\left(\frac{\pi_k}{\pi_K}\right) - \frac{1}{2\sigma^2} \{(x - \mu_k)^2 - (x - \mu_K)^2\} \\
&= \log\left(\frac{\pi_k}{\pi_K}\right) - \frac{1}{2\sigma^2} \{2(\mu_K - \mu_k)x - (\mu_K^2 - \mu_k^2)\} \\
&= \log\left(\frac{\pi_k}{\pi_K}\right) + \frac{\mu_K^2 - \mu_k^2}{2\sigma^2} + \frac{(\mu_k - \mu_K)x}{\sigma^2} \\
&= a_k + b_k x \\
\therefore a_k &= \log\left(\frac{\pi_k}{\pi_K}\right) + \frac{\mu_K^2 - \mu_k^2}{2\sigma^2}, \quad b_k = \frac{\mu_k - \mu_K}{\sigma^2}
\end{aligned}$$

#12.

(a) $\log\left(\frac{\widehat{\Pr}(Y=orange|X)}{1-\widehat{\Pr}(Y=orange|X)}\right) = \widehat{\beta}_0 + \widehat{\beta}_1 X$. Therefore, log odds of orange versus apple in my model is $\widehat{\beta}_0 + \widehat{\beta}_1 X$

(b) $\log\left(\frac{\widehat{\Pr}(Y=orange|X)}{1-\widehat{\Pr}(Y=orange|X)}\right) = (\hat{\alpha}_{orange0} - \hat{\alpha}_{apple0}) + (\hat{\alpha}_{orange1} - \hat{\alpha}_{apple1})X$. Therefore, log odds of orange versus apple in friend's model is $(\hat{\alpha}_{orange0} - \hat{\alpha}_{apple0}) + (\hat{\alpha}_{orange1} - \hat{\alpha}_{apple1})X$

(c) $\begin{cases} \hat{\alpha}_{orange0} - \hat{\alpha}_{apple0} = 2 \\ \hat{\alpha}_{orange1} - \hat{\alpha}_{apple1} = -1 \end{cases}$. However, finding a solution of 4 parameters with 2 equations is impossible.

(d) $\begin{cases} \hat{\alpha}_{orange0} - \hat{\alpha}_{apple0} = 1.2 - 3 = -1.8 = \widehat{\beta}_0 \\ \hat{\alpha}_{orange1} - \hat{\alpha}_{apple1} = -2 - 0.6 = -2.6 = \widehat{\beta}_1 \end{cases}$

(e) The models are identical with different parameterization so they should perfectly agree.

#14.

(a)

```
[92] med = df['mpg'].median()

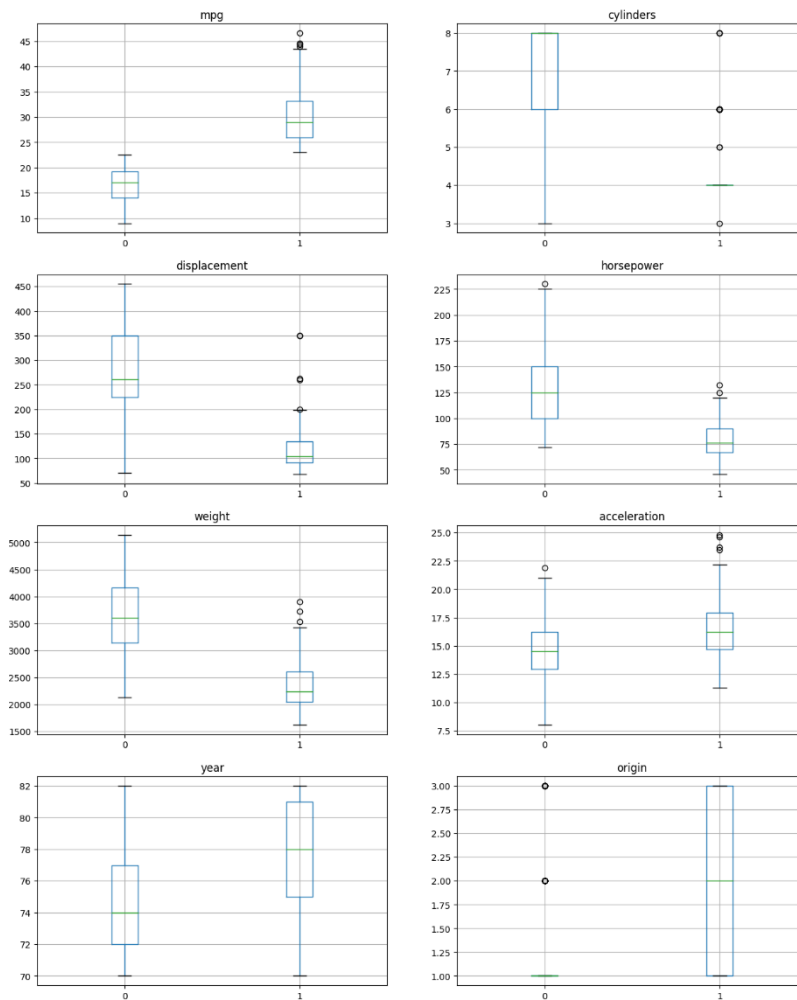
df['mpg01'] = df['mpg'].apply(lambda x: 1 if x > med else 0)
```

(b)



[Scatterplots between features]

Boxplot grouped by mpg01



[Boxplots grouped by mpg01]

'cylinders', 'displacement', 'horsepower', 'weight', 'acceleration', and 'year' seems well to fit mpg01

(c)

```
[98] x = np.array(df[['cylinders', 'displacement', 'horsepower', 'weight', 'acceleration', 'year']])
      y = np.array(df['mpg01'])
```

```
[99] # i.
      x_train, x_test, y_train, y_test = train_test_split(
          x, y, test_size=0.33, random_state=1
      )
      print(len(x_train))
      print(len(x_test))
      print(len(y_train))
      print(len(y_test))
```

```
262
130
262
130
```

※ From (d) to (h), 'cylinders', 'displacement', 'horsepower', 'weight', 'acceleration', and 'year' will be used as predictors of 'mpg01'.

(d)

```
✓ [101] lda = LinearDiscriminantAnalysis()  
0主 lda.fit(x_train, y_train)  
  
print('Test Error Rate is {0}'.format(1-lda.score(x_test, y_test)))  
  
Test Error Rate is 0.08461538461538465
```

```
✓ [102] display(confusion_table(confusion_matrix(y_test, lda.predict(x_test))))  
0主
```

	y_pred=0	y_pred=1	Total
y=0	58	10	68
y=1	1	61	62
Total	59	71	

Test error (LDA): 0.084615

(e)

```
✓ [103] qda = QuadraticDiscriminantAnalysis()  
0主 qda.fit(x_train, y_train)  
  
print('Test Error Rate is {0}'.format(1-qda.score(x_test, y_test)))  
  
Test Error Rate is 0.0692307692307692
```

```
✓ [104] display(confusion_table(confusion_matrix(y_test, qda.predict(x_test))))  
0主
```

	y_pred=0	y_pred=1	Total
y=0	62	6	68
y=1	3	59	62
Total	65	65	

Test error (QDA): 0.069230

(f)

```
✓ [105] lgr = LogisticRegression(max_iter=5000)
      lgr.fit(x_train, y_train)
```

```
print('Test Error Rate is {0}'.format(1 - lgr.score(x_test, y_test)))
```

Test Error Rate is 0.06153846153846154

```
✓ [106] display(confusion_table(confusion_matrix(y_test, lgr.predict(x_test))))
```

	y_pred=0	y_pred=1	Total
y=0	63	5	68
y=1	3	59	62
Total	66	64	

Test error (Logistic Regression): 0.061538

(g)

```
✓ [107] gnb = GaussianNB()
      gnb.fit(x_train, y_train)
```

```
print('Test Error Rate is {0}'.format(1 - gnb.score(x_test, y_test)))
```

Test Error Rate is 0.07692307692307687

```
✓ [108] display(confusion_table(confusion_matrix(y_test, gnb.predict(x_test))))
```

	y_pred=0	y_pred=1	Total
y=0	62	6	68
y=1	4	58	62
Total	66	64	

Test error (Naïve Bayes): 0.076923

(h)

```
✓ 0主 [109] testErrList = []  
      for k in range(1, 21):  
          knn = KNeighborsClassifier(n_neighbors=k)  
          knn.fit(x_train, y_train)  
  
          print('Test Error Rate(k={0}) is {1}'.format(k, 1 - knn.score(x_test, y_test)))  
          testErrList.append(1 - knn.score(x_test, y_test))  
  
Test Error Rate(k=1) is 0.10769230769230764  
Test Error Rate(k=2) is 0.13076923076923075  
Test Error Rate(k=3) is 0.13076923076923075  
Test Error Rate(k=4) is 0.11538461538461542  
Test Error Rate(k=5) is 0.13076923076923075  
Test Error Rate(k=6) is 0.13076923076923075  
Test Error Rate(k=7) is 0.09999999999999998  
Test Error Rate(k=8) is 0.08461538461538465  
Test Error Rate(k=9) is 0.09999999999999998  
Test Error Rate(k=10) is 0.09999999999999998  
Test Error Rate(k=11) is 0.12307692307692308  
Test Error Rate(k=12) is 0.11538461538461542  
Test Error Rate(k=13) is 0.12307692307692308  
Test Error Rate(k=14) is 0.11538461538461542  
Test Error Rate(k=15) is 0.11538461538461542  
Test Error Rate(k=16) is 0.11538461538461542  
Test Error Rate(k=17) is 0.12307692307692308  
Test Error Rate(k=18) is 0.08461538461538465  
Test Error Rate(k=19) is 0.10769230769230764  
Test Error Rate(k=20) is 0.09999999999999998
```

When we consider only overall test error rate, it seems **k=8** perform the best on the data set.

Chapter 5

#1.

$$f(\alpha) = \text{Var}(\alpha X + (1 - \alpha)Y) = \alpha^2 \text{Var}(X) + (1 - \alpha)^2 \text{Var}(Y) + 2\alpha(1 - \alpha)\text{Cov}(X, Y)$$

To find α that minimizes $f(\alpha)$, let's differentiate it with regard to α .

$$\frac{d}{d\alpha} f(\alpha) = 2\alpha \text{Var}(X) - 2(1 - \alpha) \text{Var}(Y) + 2(1 - 2\alpha) \text{Cov}(X, Y) = 0$$

$$\therefore \alpha = \frac{2\text{Var}(Y) - 2\text{Cov}(X, Y)}{2\text{Var}(Y) + 2\text{Var}(X) - 4\text{Cov}(X, Y)} = \frac{\sigma_Y^2 - \sigma_{XY}}{\sigma_X^2 + \sigma_Y^2 - 2\sigma_{XY}} \quad (5.6)$$

Examining $f(x)$ near $x = \alpha$ tells that $f(x)$ is minimized at $x = \alpha$

#2.

(a) $1 - \frac{1}{n}$

(b) As each bootstrap sample is a random sample, this probability is the same. $1 - \frac{1}{n}$


(c) As every sample is independent, it is $\left(1 - \frac{1}{n}\right)^n$

(d) It is $1 - \left(1 - \frac{1}{5}\right)^5 = 0.67232$

(e) It is $1 - \left(1 - \frac{1}{100}\right)^{100} = 0.63397$

(f) It is $1 - \left(1 - \frac{1}{10000}\right)^{10000} = 0.63212$

(g) The probability likely to converge in $1 - e^{-1}$ as j increases, since $\lim_{j \rightarrow \infty} \left(1 - \frac{1}{j}\right)^j = \frac{1}{e}$.



(h) 0.6398639863986398, result from bootstrapping resembles theoretical probability.

```
store = []
for i in np.arange(1, 10000):
    store += [np.sum((np.random.randint(low=1, high=101, size=100) == 4)) > 0]

np.mean(store)

0.6398639863986398
```


#5.

(a)

```
✓ [157] x = np.array(df[['balance', 'income']])  
0主 y = np.array(df.default)  
  
✓ [159] logit = LogisticRegression()  
0主 logit.fit(x, y)  
  
print(logit.intercept_, logit.coef_)  
  
[-11.54047812] [[5.64710797e-03 2.08091985e-05]]
```

(b)

i.

```
[118] # i.  
x_train, x_test, y_train, y_test = train_test_split(  
    x, y, test_size=0.33, random_state=1  
)  
print(len(x_train))  
print(len(x_test))  
print(len(y_train))  
print(len(y_test))  
  
6700  
3300  
6700  
3300
```

ii.

```
✓ [161] # ii.  
0主 logit = LogisticRegression()  
logit.fit(x_train, y_train)  
  
print(logit.intercept_, logit.coef_)  
  
[-1.18250162e-06] [[ 0.00039328 -0.00012161]]
```

iii.

```
[121] # iii.  
  
probs = logit.predict_proba(x_test)  
predict = (probs[:, 1] > 0.5).astype(int)  
result = ['Yes' if item==1 else 'No' for item in predict]  
  
[122] display(confusion_table(confusion_matrix(y_test, result)))
```

	y_pred=0	y_pred=1	Total
y=0	3200	0	3200
y=1	100	0	100
Total	3300	0	

iv.

Validation set overall error rate is $100/3300 = 0.030303$

(c)

random_state: 2

	y_pred=0	y_pred=1	Total
y=0	3212	2	3214
y=1	86	0	86
Total	3298	2	

[Validation Set Overall Error Rate]

- random_state (2) => $(86+2)/3300 = 0.026667$

random_state: 3

	y_pred=0	y_pred=1	Total
y=0	3179	8	3187
y=1	74	39	113
Total	3253	47	

- random_state (3) => $(74+8)/3300 = 0.024848$

random_state: 4

	y_pred=0	y_pred=1	Total
y=0	3180	12	3192
y=1	69	39	108
Total	3249	51	

- random_state (4) => $(69+12)/3300 = 0.024545$

The results obtained are variable and depend on the samples allocated to training vs. test.

(d)

random_state: 2

	Pred: No	Pred: Yes	Total
No	3212	2	3214
Yes	86	0	86
Total	3298	2	

[Validation Set Overall Error Rate (Including 'Student')]

- random_state (2) => $(86+2)/3300 = 0.026667$

random_state: 3

	Pred: No	Pred: Yes	Total
No	3173	14	3187
Yes	91	22	113
Total	3264	36	

- random_state (3) => $(91+14)/3300 = 0.031818$

random_state: 4

	Pred: No	Pred: Yes	Total
No	3176	16	3192
Yes	87	21	108
Total	3263	37	

- random_state (4) => $(87+16)/3300 = 0.031212$

Including student does not seem to make a substantial improvement to the test error.

#6.

(a)

```
smf.logit('default ~ income + balance', data=temp_df).fit().summary()
```

Optimization terminated successfully.
Current function value: 0.078948
Iterations 10

Logit Regression Results

Dep. Variable:	default	No. Observations:	10000
Model:	Logit	Df Residuals:	9997
Method:	MLE	Df Model:	2
Date:	Fri, 29 Sep 2023	Pseudo R-squ.:	0.4594
Time:	08:34:34	Log-Likelihood:	-789.48
converged:	True	LL-Null:	-1460.3
Covariance Type:	nonrobust	LLR p-value:	4.541e-292

	coef	std err	z	P> z	[0.025	0.975]
Intercept	-11.5405	0.435	-26.544	0.000	-12.393	-10.688
income	2.081e-05	4.99e-06	4.174	0.000	1.1e-05	3.06e-05
balance	0.0056	0.000	24.835	0.000	0.005	0.006

Possibly complete quasi-separation: A fraction 0.14 of observations can be perfectly predicted. This might indicate that there is complete quasi-separation. In this case some parameters will not be identified.

standard error of estimated coefficient(income): 4.99e-06

standard error of estimated coefficient(balance): 0.000

(b)

```
[298] def boot_fn(default):  
      mod1 = smf.logit('default ~ income + balance', data=default).fit(maxiter=10000)  
      coef_income = mod1.params[1]  
      coef_balance = mod1.params[2]  
      return [coef_income, coef_balance]
```

```
[299] boot_fn(temp_df)  
  
Optimization terminated successfully.  
Current function value: 0.078948  
Iterations 10  
[2.0808975528986992e-05, 0.005647102950316493]
```

(c)

```
✓ [300] def bootstrap_sample(df, num_samples=None):  
0主      n = df.shape[0]  
  
      if num_samples == None:  
          num_samples = df.shape[0]  
  
      bootstrap_samples = np.random.choice(n, num_samples, replace=True)  
      bootstrap_df = df.iloc[bootstrap_samples, :]  
      return bootstrap_df  
  
✓ [301] coef_list = []  
5主      for _ in tqdm(range(100)):  
          try:  
              sample_coefs = boot_fn(bootstrap_sample(temp_df))  
              coef_list.append(sample_coefs)  
          except Exception as e:  
              continue
```

(d)

```
✓ [302] pd.DataFrame(coef_list, columns=['income', 'balance']).describe()  
0主
```

	income	balance
count	100.000000	100.000000
mean	0.000020	0.005652
std	0.000004	0.000225
min	0.000010	0.005201
25%	0.000018	0.005506
50%	0.000021	0.005649
75%	0.000023	0.005800
max	0.000031	0.006193

From (a), estimated coefficients and standard errors of these coefficients for logistic regression of income and balance was $[2.081e-05, 0.0056]$, and $[4.99e-06, 0.000]$ respectively. According to the bootstrapping method, we obtained coefficients $[0.000020, 0.005639]$, standard errors $[0.000006, 0.000228]$ for these coefficients. From this result, one can conclude that result from bootstrapping resembles well with the original one.

#7.

(a)

```
[308] # make dummy variables
temp_df = df.copy()
temp_df['Direction_Up'] = temp_df['Direction'].apply(lambda x: 1 if x=='Up' else 0)

[309] logit = LogisticRegression()
logit.fit(np.array(temp_df[['Lag1', 'Lag2']]), np.array(temp_df['Direction_Up']))

LogisticRegression
LogisticRegression()

print(logit.intercept_, logit.coef_)

[0.22122502] [[-0.03869814  0.06020749]]
```

estimated coefficient (Lag1): -0.03869814

estimated coefficient (Lag2): 0.06020749

(b)

```
[311] logit = LogisticRegression()

# predicts Direction Using Lag1 and Lag2 using all but the first observation
x_train = np.array(temp_df.loc[1:, ['Lag1', 'Lag2']])
y_train = np.array(temp_df.loc[1:, 'Direction'])

x_test = np.array(temp_df.loc[0, ['Lag1', 'Lag2']])
y_test = np.array(temp_df.loc[0, 'Direction'])

logit.fit(x_train, y_train)
print(logit.intercept_, logit.coef_)

[0.22324404] [[-0.03840931  0.06080633]]
```

estimated coefficient (Lag1): -0.03840931

estimated coefficient (Lag2): 0.06080633

(c)

```
✓ [335] y_test
0主
array('Down', dtype='<U4')

✓ [336]
0主
prob = logit.predict_proba(x_test.reshape(-1, 2))
predict = (prob[:, 1] > 0.5)
print('Up' if predict[0] == True else 'Down'))

Up
```

It isn't correctly classified.

(d)

```
✓ [348]
12 主
loocv = LeaveOneOut()
real_y_list = []
pred_y_list = []
errors = np.zeros(temp_df.shape[0])

for i, (train_index, test_index) in enumerate(loocv.split(np.array(temp_df))):
    # i.
    train_data = temp_df.iloc[train_index, :]
    test_datum = temp_df.iloc[test_index, :]

    x_train = np.array(train_data[['Lag1', 'Lag2']])
    y_train = np.array(train_data['Direction'])
    x_test = np.array(test_datum[['Lag1', 'Lag2']])
    y_test = np.array(test_datum['Direction'])

    logit = LogisticRegression()
    logit.fit(x_train, y_train)
    # ii.
    prob = logit.predict_proba(x_test)

    # iii.
    predict = (prob[:, 1] > 0.5).astype(int)
    predict = ['Up' if item==1 else 'Down' for item in predict]

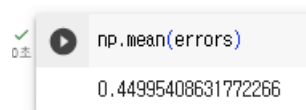
    # iv.
    if (predict[0] != y_test.reshape(-1, 1)[0][0]):
        errors[i] = 1

    real_y_list.append(y_test[0])
    pred_y_list.append(predict[0])

✓ [349]
0主
display(confusion_table(confusion_matrix(real_y_list, pred_y_list)))
```

	Pred: Down	Pred: Up	Total
Down	34	450	484
Up	40	565	605
Total	74	1015	

(e)



```
np.mean(errors)
```

0.44995408631772266

The LOOCV test error rate is 45% which implies that our predictions are marginally more often correct than not.